

1 Names and Emails

- Aleksandar Makelov — amakelov@college.harvard.edu
- Ben Wetherfield — bwetherfield@college.harvard.edu
- Chan Kang — chankang@college.harvard.edu
- Michael Fountaine — mfount@college.harvard.edu

2 Overview

Problem: Using Coq, verify Timsort, python’s preferred sorting algorithm! **[Improve motivation? Possibly include discussion of failures of implementations of Timsort in languages other than Python 3.x.]**

Solution sketch: We will take an incremental approach: We’re going to start with an abstract "implementation" of Timsort, possibly black-boxing certain difficult-to-implement features such as heuristics used in the algorithm, and implement a verified version of that implementation to sort lists of natural numbers (defined inductively). Timsort is a hybrid sorting algorithm, requiring other sorting sub-algorithms and some simple data structures, so we will divide our project into interfaces for each of these, along with interfaces for proof tactics. We’re going to hone our sorting algorithm verification chops by working through insertion sort and mergesort, both of which are part of Timsort, followed by tree sort and heap sort to test our data structures, followed by Timsort itself!

Goals: Primarily, we’d like to verify Timsort as a way to learn more about Coq and certified programming. (If Timsort proves to be too complicated (by our estimation by the time the final spec is due), we will be able to accomplish this same learning outcome by working through insertion, merge, tree, and heap sorts. Or, perhaps, we will include another hybrid search algorithm instead of Timsort.)

3 Prioritized Feature List

Note: Generally, we need to read more about Coq to understand where to draw the line between “core” and “cool”; that is, how long these features will take to implement, including proofs.

Additionally, we are unsure of whether to verify the in-place, imperative versions of these algorithms or the pure versions. The deciding factor will most likely be the difficulty of implementing arrays (and references) in Coq.

Core Features

- **Fundamentals.** Ordered lists, natural numbers (defined inductively), lists and arrays, stacks (to be used as Timsort’s memory and temporary memory).
- **Insertion sort.** Verified insertion sort of lists of natural numbers.
- **Merge sort.** Verified merge sort of lists of natural numbers.
- **Trees.** Binary search trees of natural numbers.
- **Tree Sort.** Verified tree sort of binary search trees of natural numbers.
- **Heaps.** Priority queues of natural numbers.
- **Heap Sort.** Verified heap sort of binary search trees of natural numbers.

- **Simplified Timsort.** This will operate on lists of natural numbers, represented perhaps as trees or priority queues.

Timsort is a hybrid of insertion sort and mergesort, plus some heuristics about memory management and other optimizations. Our “simplified timsort” will omit these heuristics when possible.

Cool Extensions

- **Timsort.** Verified timsort; that is, simplified timsort plus heuristics. This version would have the same time asymptotics as python’s implementation of Timsort: $\Theta(n)$ best case, $\Theta(n \log n)$ average case, $\Theta(n \log n)$ worst case, where n is the length of the list.
- **Polymorphic timsort.** A verified timsort that works on polymorphic lists. (Unlike in OCaml, this seems actually to be a good bit of work in Coq, having to prove things about each subtype of our ordered set type.)
- **Export to OCaml.** One of Coq’s initial purposes is to export verified OCaml (or Haskell or Scheme) code; for this cool extension, we can try to turn some or all of our verified algorithm implementations into usable OCaml.
- **Python’s timsort.** This final possible extension could be to modify our verified Timsort algorithm to use exactly the heuristics found in the current Python 3.x release, as opposed to the heuristics we end up using in our implemented version of the full Timsort algorithm. We could also attempt to use the same space complexity of Timsort, $O(n)$.

4 Technical Specification

Data Structures We will, of course, prove all methods associated with these types/structures.

- Ordered set.
- Natural numbers.
- List of natural numbers.
- Array of natural numbers. (Unclear if necessary.)
- Stack.
- Tree.
- Heap (a.k.a. Priority Queue).

Algorithms These are the algorithms and subroutines we’ll need for the mutable (and often in-place) versions of the sorts, which we expect to be harder; if we decide to do the pure versions, we will make adjustments accordingly. The ideas are based on the paper [3], which describes how to prove correctness for insertion sort, quicksort and mergesort. It also provides some general predicates for formally verifying sorting algorithms, like `(sorted t)` which tests for sortedness, and `(permutation t t’)` which tests whether t' is a permutation of t .

- **Insertion sort:** The subroutine `(insertion n t)` takes as input an array t where the first $n - 1$ elements are sorted, and returns an array where the first n elements are sorted, by inserting the n -th element in its place. Then insertion sort itself proceeds by running `(insertion n t)` for $n = 1, 2, 3, \dots, N$ for a list t of length N .

- **Mergesort:** Here the basic subroutine is `(merge t)` which takes a list t in which the first and second halves are sorted, and merges them. Here we'll need to use some extra memory. Then mergesort proceeds by calling itself recursively on each half, and then merging them together.
- **Tree sort:** We first build a binary search tree from the array by a subroutine `(bst_build t)`; then we have another subroutine `(bst_traverse t)` that traverses it in-order.
- **Heap sort:** The algorithm proceeds by making the array into a heap in-place, and then gradually pushing the largest elements to the right side of the heap. Here, we're representing a binary tree implicitly as an array, where the children of the i -th element are the $2i + 1$ -th and $2i + 2$ -th elements. We have the following subroutines:
 - Predicate `(heap t n k)` checks if in a list t , the tree of elements of index $\leq n$ rooted at the k -th element.
 - Predicate `(inf tree t n v k)` checks if in a list t , every element in the tree of elements of index $\leq n$ rooted at the k -th element is less than v .
 - Subroutine `(downheap t k n)` takes a list t where the tree of elements of index $\leq n$ rooted at the k -th element is a heap *except* possibly for the root node. It then makes a bunch of swaps that make this tree into a heap.

Having these subroutines, heapsort is easy: first build the heap using `downheap` a bunch of times, and then swap the 1-st element with the i -th and rebuild the heap up to index $i - 1$, for $i = N, N - 1, \dots, 1$

- **Simple timsort:** The idea of timsort is to use the sorted chunks (called *runs*) often present in real-world data. At a high level, the algorithm works by first making sure all runs are of some minimum length (using insertion sort if necessary), and then passing over the list and merging the runs in an intelligent way.

As we do the second pass, we push the starting index and length of each run on a stack, and merge consecutive runs in order to preserve a certain invariant of the stack: if A, B, C are the lengths of three consecutive runs on the stack (with A being the topmost), we require that $B > C$ and $A > B + C$.

So at any point we end up with a bunch of runs whose sizes grow faster than the Fibonacci sequence, i.e. at least exponentially fast; so it's easy to see there are at most logarithmically many runs at each point in time, which is important for memory reasons and for running time (it's faster to merge a sequence of lists with exponentially-increasing size than to naively merge a list split into equal parts).

The way the invariant is preserved is where things get weird: apparently, there is a bug in the algorithm on the wikipedia page (which attempts to restore the invariant by only fixing the top three runs), but there is a fix [4] we have to look into.

It's hard to know in advance what interface would be best for our proof, but here's a guess:

- Predicate `(run t i j)` - this checks if the subarray between the i -th and j -th index of t is a run, i.e. is already in sorted order.
- Subroutine `insertion-sort` - this is just as insertion sort above.
- Subroutine `merge` - this is supposed to take the next run and do any merges necessary to preserve the invariant; of course, it will be based on the merge procedure.
- Now piecing together all the parts should be easy. In the end, we end up with a sequence of runs that we merge in the obvious way.
- **Timsort:** Additionally, there are some memory optimizations ('galloping') during the merges that we might try to include if we end up having the time for it.

Proofs *Note:* We will be able to better characterize proofs after working through *Software Foundations* [2] (See “Next Steps” below.)

At this stage, we know:

- Because Coq proofs are very formal, they suffer from an inability to convey human-to-human understanding that is afforded by semi-formal, traditional mathematical proof. To aid debugging and authoring of proofs in Coq, *SF* [2] recommends two strategies, “Informal Proof” and “Human Proof”, that by encouraging verbosity and a high degree of refactoring help to restore the “understanding” aspect of a proof.
- Proofs in coq are guided by tactics. These are like hints you give to the prover, such as “rewrite this as the other side of this equality we already proved”, “rewrite both sides in a canonical way”, etc.
- We expect to use the “Reflexivity” keyword frequently for verifying more basic properties of our data structures. Proofs in Coq that use reflexivity as a tactic are generally very simple proofs using, say, rewrite rules.
- When we get onto proofs of correctness of algorithms, we will have to break proofs into subcomponents.
- There are various ways of doing this:
 1. Example, Theorem, Lemma, Fact and Remark keywords (all functionally the same) break up larger proof into subproofs.
 2. The Case proof tactic (chapter 3 of *Software Foundations*) can be used to break a proof clearly into cases (as in the familiar structure of an induction proof).
 3. This could break a proof into say a case of the Reflexivity tactic and the Induction tactic. Our reading of the references and some example Coq code indicates that this is a common formulation that appears when doing structural induction.
 4. Ltac can be used to create common tactics.

5 Next Steps

[Revise these!]

- Install Coq 8.4-pl5 and ensure that it’s running correctly. Also install an IDE, most likely emacs with Proof General.
- Read early parts of *SF* [2] and *CPDT* [1], doing exercises to familiarize ourselves with Gallina (Coq’s functional language, similar to Caml) syntax and Coq’s Proof functionality.
Specifically, we will initially work through at least the 1-star and 2-star exercises in the opening three chapters of *SF*. (This will give us the implementations of lists and nats.)
- In a similar vein, get an idea of how long it would take to verify correctness of the basic sorting algorithms
- Understand timsort in more depth.

References

- [1] Chlipala, Adam. *Certified Programming with Dependent Types*.
- [2] Pierce, Benjamin, et al. *Software Foundations*.

- [3] Filliatre, Jean-Christophe and Magaud, Nicolas. *Certification of sorting algorithms in the system COQ*
- [4] de Gouw, Stijn and Rot, Jurriaan and de Boer, Frank S and Bubel, Richard and Hahnle, Reiner. *OpenJDK's java. utils. Collection. sort () is broken: The good, the bad and the worst case*