

1 Overview

Problem: Using Coq, verify Timsort, python's preferred sorting algorithm and widely used in a variety of languages! Yet implementations of Timsort have been known to fail. We turn to Coq for its support of *verifiable* code and draw on its formal proof techniques, which assures us that we won't run into nasty surprises, dreaded holes in the correctness of the algorithm! Coq supports the exporting of programs into other languages such as ocaml, so we expect to render a sorting procedure that is not only water-tightly verified, but also usable in a broader context. Timsort is a hybrid of insertion sort and mergesort, plus some heuristics about memory management and other optimizations. Our simplified version will include a reduced version of these heuristics.

Solution sketch: We will take an incremental approach (and all of our algorithms will be functional, using persistent data structures). We're going to start with a simplified version of Timsort, hybridizing mergesort and insertion sort with a small subset of the heuristics used by Timsort in full. At first, all of these components of this simplified Timsort will be independently verified; the combined algorithm implementation will be thoroughly unit-tested. For short, call this algorithm Simsort. Implementing Simsort will be the conclusion of our core functionality.

Next, we will implement heaps and a verified heapsort; by replacing insertion sort with heapsort in Simsort, we should get a constant-factor time improvement. Then, our primary goal beyond core functionality will be verification of Simsort. From there, we will implement extra extensions, as discussed below, possibly adding more heuristics to Simsort, approaching verification of Timsort in full.

Goals: Primarily, we'd like to verify Timsort (i.e., Simsort) as a way to learn more about Coq and certified programming.

2 Prioritized Feature List

Note: All algorithms and data structures used in this project will be functional; in particular, we'll use persistent data structures.

Core Features

- **Fundamentals.** Booleans, natural numbers (defined inductively), polymorphic lists, stacks (for very basic representations of memory needed within Timsort heuristics).
- **Verified insertion sort.** Verified insertion sort of lists of natural numbers.
- **Verified merge sort.** Verified merge sort of lists of natural numbers.
- **Simsort.** Fully tested implementation of Simsort (our hybridization of verified merge sort, verified insertion sort, and a modified subset of the heuristics used in Timsort). We still need to determine exactly which subset of heuristics we will use and how we will modify them.

Cool Extensions

- **Heaps.** Polymorphic priority queues. If it provides an advantage in asymptotics, we will use heaps to re-implement stacks.
- **Heap Sort.** Verified heap sort of heaps of natural numbers. This will operate on lists of natural numbers, represented perhaps as trees or priority queues.
- **Augmenting Simsort with heapsort.** Fully tested implementation of Simsort, with heapsort replacing insertion sort, for a slight improvement in asymptotics.

- **Verified Simsort.** This is our main goal beyond core functionality. We will improve the fully tested Simsort to a rigorously verified Simsort (using Coq).
- **Passing foreign tests.** We have come across a few known to be broken implementations of Timsort in certain languages (e.g., Java’s clone of Timsort, early versions of Python 2.x’s Timsort). For this cool extension, we would take some of the failing test cases for those other implementations, adapt them to use the same heuristic assumptions that we’ve used with Simsort, and show that our verified Simsort passes those tests.
- **Adding more heuristics.** If we make it this far, we will add more heuristics to Simsort, showing that each addition passes verification and doesn’t break invariants, working our way gradually to a verified, functional Timsort in full.

3 Technical Specification

3.1 Interfaces.

References to code written thus far are included throughout this section. `chicken` is the root directory of our github repository, located at <https://github.com/mfount/chicken>.

Also, all libraries originating from SF [2] (which we’ve populated with exercise solutions), will eventually be refactored into neater, application-specific libraries once we have a clearer picture of how we will refactor code to optimize proofs.

Data Structures, related methods, and proofs of them.

- **Booleans.** Our code for booleans so far is contained in `chicken/coq/Basics.v`, a library provided with SF [2], which we have filled out with solutions to the relevant exercises. This includes axioms and definitions pertaining to booleans (e.g. `negb false = true`), proofs of simple lemmas (primarily done using the `destruct` tactic to prove by case exhaustion) familiar binary operations, and proofs of them.
- **Natural numbers.** There is code for nats in `chicken/coq/Basics.v` and `chicken/coq/Induction.v`, another SF library that we have supplemented with solutions to relevant exercises. Nats are defined inductively using a unary representation (as Peano naturals). Most of the lemmas pertaining to nats are proven using the `induction` tactic.

For our purposes, the set of natural numbers is most usefully structured as an ordered set, so we’ve verified (or made progress toward verifying) the order properties of the natural numbers as comparison methods, for example `ble_nat` (`<=` for nats) as well as `blt_nat` (`<` for nats).

- **Polymorphic lists.** We will be sorting objects of the type `list nat` (informally, “natlists”) throughout this project. Also, we will use lists of lists, `list (list nat)` in Simsort. Definitions of polymorphic lists are found in `chicken/coq/Poly.v` and `chicken/coq/Lists.v` other SF libraries modified as described above. (The former defines polymorphic lists, and the latter defines, specifically, lists of nats.)
- **Stack.** We’re going to implement this using polymorphic lists. Methods will include `push` (implemented as the familiar list `Cons` construction) and `pop` (implemented as `hd_opt`). Proofs of consistency of these will be based on existing exercises in SF, which we have already done in working through the textbook. (See `chicken/coq/Poly.v`.)
- **Heap.** See `chicken/coq/Heap.ml` for a standard priority queue interface. We will transliterate this to Coq (technically, to Gallina, Coq’s functional language subset). Also, we expect many proofs of the methods used within to resemble the lemma and tactic structure of polymorphic lists.

Algorithms and proofs of them. In this section, we outline the methods required by each algorithm and the necessary lemmas that we will need to implement to verify these methods. Additionally, we will need pre- and post-conditions for the algorithm as a whole, but what exactly will constitute these should become apparent as we make progress over this first week.

- **Fundamental methods and proofs relevant to all sorting algorithms.**

- **Mergesort.**

Mergesort:

- Procedure ('merge t s'):

takes input lists 't', 's' that are sorted

returns an array that is 's' and 't' merged together (and sorted)

- Merge sort:

takes input a list 't'

calls itself recursively on each half of the list, then calls 'merge'.

- **Insertion sort.**

- Procedure ('insert n t'):

takes input an list 't' where the last 'n-1' elements are sorted

returns list 't' where last 'n' elts are sorted, by inserting 1st element

- Insertion sort:

takes input a list 't_1'

run ('insert 0 t_{n-1}'), ('insert 1 t_{n-2}'), ..., ('insert {n-1} t_{0}')

where 't_i' is the result of ('insertion (n-i) t_{i-1}')

Proof: for the proof we prove insert and then InsertionSort on top of that. We will need IsSorted, which we prove elsewhere.

See 'chicken/coq/InsertionSort.v'.

- Procedure ('insert n t'):

takes input an list 't' where the last 'n-1' elements are sorted

returns list 't' where last 'n' elts are sorted, by inserting 1st element

- Insertion sort:

takes input a list 't_1'

run ('insert 0 t_{n-1}'), ('insert 1 t_{n-2}'), ..., ('insert {n-1} t_{0}')

where 't_i' is the result of ('insertion (n-i) t_{i-1}')

Proof: for the proof we prove insert and then InsertionSort on top of that. We will need IsSorted, which we prove elsewhere.

See 'chicken/coq/InsertionSort.v'.

- **Heap sort.**

Heapsort:

- Procedure ('heapify t'):

takes input a list 't'

returns a heap data structure with elements the element of the list

- Procedure (`heapsort_step h t`):
 takes as input a heap (a min-heap) `h` and a list `t`
 returns a new heap `h'` and new list `t'` where we moved the min element
 from the heap `h` to the beginning of `t`.

- Heap sort:
 takes as input a list `t`
 runs `heapify` on `t`
 then runs `heapsort_step` until the heap is empty.

- **Simsort.**

Simple timsort (`simsort`):
 - Procedure (`runs t`):
 takes as input a list `t`
 returns: consider a permutation of `t` where consecutive runs have
 length at least 64 (or some other constant, we'll see)
 then (`runs t`) returns a list of these runs

Uses either insertion sort (original timsort) or heapsort (our
 modification) to sort the runs when necessary

- Procedure (`merge_runs T`)
 takes as input a list of lists `T`, as the output of (`runs t`).
 returns a list of lists that represents the result of merging the runs
 following the timsort heuristic:
 the runs are put on a stack, and then if X, Y, Z are the lengths of the top 3
 runs, the algorithm merges the runs
 until the invariant $X > Y + Z, Y > Z$ is satisfied.

stack can be done as a list or using dictionary type which is defined in
`'chicken/coq/Poly.v'`.

4 Timeline

As of Friday 17 April at 5pm, we have exactly two weeks to finish our project.

Week One.

- Isolate axiom definitions from the first 4 chapters of Software Foundations that we will use. Get them working as a set of libraries that can be Required.
- As a group, do all of these exercises (proof of the consistency of the axiom definitions that we need). As individuals, learn the contents of the first 5 chapters.
- Complete proofs of `is_sorted_le` in `chicken/coq/IsSorted.v`. Complete proofs of `insertion_sort` in `chicken/coq/InsertSort.v`. (Both of these libraries are ours, not originating from SF, and already exist in part.)

- Do the same for `merge_sort`.
- Get an implementation of heaps working (in line with the pseudo-code signature given in the final spec) and prove half of the operations.

Week Two.

- Complete proof of the correctness of operations of heap.
- Implement `heap_sort` and prove it. Easy after the work we have done for heaps.
- combine merge and insert into `simsort`. Export `simsort` to OCaml and write test suite.
- Revise `simsort` to implement heaps; export to OCaml; ensure that the test suite is still passed.
- Verify `simsort`.
- Further tasks if possible. (See “Cool Extensions”, as outlined above.)

5 Detailed Progress

All code written so far is contained in the directory `chicken/coq`, where, again `chicken` is the root directory of our github repository, located at <https://github.com/mfount/chicken>.

Here is an overview of each of the files contained within (with the exception of those files generated by the compiler):

1. `#Data_structure.v#`
2. `Basics.glob`
3. `Basics.html`
4. `Basics.v`
5. `Basics.vo`
6. `Data_structure.v`
7. `Heap.ml`
8. `Heap.v`
9. `Induction.glob`
10. `Induction.html`
11. `Induction.v`
12. `Induction.vo`
13. `InsertionSort.v`
14. `IsSorted.glob`
15. `IsSorted.v`
16. `IsSorted.vo`
17. `Lists.glob`

18. `Lists.html`

19. `Lists.v`

20. `Lists.vo`

21. `Poly.glob`

22. `Poly.v`

23. `Poly.vo`

References

[1] Chlipala, Adam. *Certified Programming with Dependent Types*.

[2] Pierce, Benjamin, et al. *Software Foundations*.